

Kapil Wanaskar

 San Francisco, CA, US  kapilw25@gmail.com  +16698379485  LinkedIn  Google Scholar

Work Experience

- First American Title, Senior ML Engineer** Apr 2025 – Present | Remote, CA, USA
- Engineered MultiAgent [WebSearch + Verification] workflow automating 250,000+ daily repetitive tasks, eliminating \$1.25M in manual processing costs (\$5 per task) while achieving 98% accuracy rate [Patent Pending]
 - Optimizing model costs by 78% through knowledge distillation: Fine-tuned Qwen2.5-VL-7B to replace Gemini-2.5-Pro-API, deployed via Ollama achieving 2.3x faster inference speeds in production MultiAgent workflow
- Amazon Web Services (AWS), Applied ML Engineer** May 2024 – Mar 2025 | Cupertino, CA, USA
- Implemented GRPO, an advanced Reinforcement Learning (RL) fine-tuning strategy, to optimize Llama3.1; improved training efficiency by 40% over PPO while preserving accuracy across 100K+ samples
 - Optimized large-scale distributed model training across 100+ Trainium accelerators with int4 quantization, implementing efficient data parallelism that reduced inter-node communication overhead by 35% and memory usage by 60%
- Intuitive Surgical, Software Engineer - ML** May 2023 – May 2024 | Sunnyvale, CA, USA
- Built FastAPI + VectorDB-based inference pipeline to detect out-of-distribution robotic surgery logs; achieved 98% precision, enabling early-stage anomaly flagging by the clinical safety team
 - Supervised fine-tuned (SFT) LLM via PEFT (LoRA) on few-shot human-labeled feedback
 - Post-trained embedding encoders and re-indexed FAISS via Online Reinforcement Learning (RL) for similarity updates, improving security by 13%
 - HyperParameter tuned 130,000+ variations of unsupervised models on 150+ GB data using SageMaker + MLflow; achieved 92% precision and 99.9% accurate training inputs
- Vectorr.in, Software Engineer** Mar 2018 – Jul 2022 | Mumbai, India
- Engineered unsupervised customer segmentation system 10k+ (daily) visits stored in Snowflake database, surging customer satisfaction from 3.1 to 4.8.
 - Integrated Apache Kafka and Superset to segment real-time audience data for digital marketing while training Unsupervised models on AWS EC2, amplifying ROI by 23%.

Research Publications

- A Framework for Switchable LLM Alignment via CITA – Contrastive Instruction-Tuned Alignment**  2026
- Achieved 86.7% instruction-alignment efficiency on Llama-3.1-8B with 25% higher reward margins (7.5 vs 6.0) **outperforming DPO** by 30.6pp, **GRPO** by 50.6pp, and **PPO** by 66.3pp across multi-dimensional evaluation
 - Created ECLIPTICA benchmark with 3,000 prompt across 10 instruction types and 5 benchmarks
 - Developed CITA algorithm with unified training pipeline (SFT→DPO→CITA) combining contrastive preference optimization and mandatory KL anchor, validated through Optuna-based hyperparameter search across 13 trials
- A Comprehensive Dataset for Human vs. GenAI Image Detection**  2026
- Created MS COCOAL dataset with 96,000 semantically-aligned real/synthetic pairs from 5 Text2Image models for controlled artifact studies
 - Established dual-task benchmark achieving 80.1% binary detection and 44.9% model attribution
 - Designed frequency-domain ResNet-50 robustness framework, detecting with 80.1% accuracy across 4 perturbations
- Multimodal Benchmarking and Recommendation of Text-to-Image Generation Models, IEEE CISOSE 2025**  2025
- Received the "BDS Best Student Paper" award
 - Evaluated 12+ (text-to-image) models (Stable Diffusion, CogView, FLUX, etc.) with ground truth from DeepFashion Multimodal dataset for alignment
 - Designed Weighted Score metric combining CLIP-Score, LPIPS, FID, MRR& Recall@3 via min-max normalization
 - Integrated metadata features and CLIP embeddings to align generated with ground truth image and prompt context
 - Metadata-augmented models (Flux, InContext LoRA) showed ~19% higher Weighted Score & ~15 point FID reduction
- Prompt Recommendations for AI art, IEEE AIKE, California, USA**  2023
- Extracted features of 5000 images via text embeddings and ensemble models
 - Proposed Graph-based evaluation of 3 recommendation Algorithms and Community Detection Algorithms, via analyzing absence of ratings or preference scores

Education

- MS in Artificial Intelligence, Computer Engineering,** CA, USA
San José State University
- Master of Computer Integrated Manufacturing and Bachelor of Engineering, Indian Institute of Technology (IIT) Bombay** Mumbai, India