Kapil Wanaskar

Cupertino, CA, USA | kapilw25@gmail.com | +16698379485 | LinkedIn | Google Scholar

Work Experience

Amazon Web Services (AWS), *LLM Optimization / Applied ML* May 2024 – present | Cupertino, CA, USA
- Enabled MoE (DBRX, Mistral) with stability-aware expert routing; improved LLM training efficiency by 18%
- Built eval pipelines for MoE variants (dropless/dropping); reducing convergence errors by 30%
- Optimized PyTorch-based NxD parallelism for MiXtral-style LLMs, enabling 12% faster training
- Developed activation/gradient drift tools (CPU vs. TRN); flagged 95% silent regressions

Intuitive Surgical, *Software (Machine Learning) Engineer* May 2023 – May 2024 | Sunnyvale, CA, USA
- Developed FastAPI-based ML inference pipeline with VectorDB to detect 98% precision "Unknown" attacks
- Supervised fine-tuned LLMs (Claude, Llama) via PEFT (LoRA) on few-shot human-labeled feedback
- Post-trained embedding encoders and re-indexed FAISS for similarity updates, improving security by 13%
- Fine-tuned 130,000+ unsupervised models on 150+ GB data using SageMaker + MLflow; achieved 92% precision and 99.9% accurate training inputs

Vectorr.in, *Software Engineer* Mar 2018 – Jul 2022 | Mumbai, India
- Clustered 10k+ (daily) visits stored in Snowflake database, surging customer satisfaction from 3.1 to 4.8.
- Integrated Apache Kafka and Superset to segment real-time audience data for digital marketing while training Unsupervised models on AWS EC2, amplifying ROI by 23%.
- Deployed Docker via CI/CD for automating deployment, achieving a 43% reduction in data overhead.

Research Publications

Multimodal Benchmarking and Recommendation of Text-to-Image Generation Models
- Evaluated 12+ (text-to-image) models (Stable Diffusion, CogView, FLUX, etc.) with ground truth from DeepFashion Multimodal dataset for alignment
- Designed Weighted Score metric combining CLIP-Score, LPIPS, FID, MRR& Recall@3 via min–max normalization
- Integrated metadata features and CLIP embeddings to align generated with ground truth image and prompt context
- Metadata-augmented models (Flux, InContext LoRA) showed ~19% higher Weighted Score and ~15 point FID reduction

Evaluation of Local LLM models for shopping recommendation
- Benchmarked "Llama 3.1," "Gemma," "phi3.5," and "Qwen" on Answer Relevancy, Contextual accuracy, etc, reducing model selection time by 25%
- Transitioned from RAG to FAISS embeddings to optimized vector retrieval, reducing hallucinations by 30%.
- Evaluated LLMs using "Nemo-Mistral" , fine-tuning benchmarking scripts to cut evaluation runtime by 40%

Prompt Recommendations for AI art, *IEEE AIKE, California, USA*
- Extracted features of 5000 images via text embeddings and ensemble models
- Proposed Graph-based evaluation of 3 recommendation Algorithms and Community Detection Algorithms, via analyzing absence of ratings or preference scores

Surveillance Drone Cloud and Intelligence Service, *IEEE (MobileCloud), Greece*
- Proposed a surveillance drone cloud for efficient utilization of cloud computing and real-time data sharing

Detection of Cyber Security Threats using IOT Deep Learning
- Suggested TensorFlow deep neural system to classify stolen programming with source code literary theft

Real World Use of Deep Learning Models for Cyber Security in IoT Network
- Presented Deep reinforcement learning models for cyber security in IoT (Internet of Things) networks

Analyzing Effect of Workpiece Stiffness Variation on the Stability in Flank Milling of an Impeller Blade
- Scrutinized FFT (fast-fourier-transform) plots, chatter boundary plots, and stability region diagrams

Education

MS in Artificial Intelligence, Computer Engineering, CA, USA
San José State University

Master of Computer Integrated Manufacturing and Bachelor of Engineering, *Indian Institute of Technology (IIT) Bombay* Mumbai, India